

Intermediate R: tidyverse

PSYC 2020-A01 / PSYC 6022-A01 | 2025-09-19 | Lab 5

Jessica Helmer

Outline

- Assignment 4 Review
- `tidyverse` Workflow
- Full Analysis

Learning objectives:

R: `tidyverse`

Housekeeping

[input any housekeeping items here]

Assignment 4 Review


[input any Assignment 4 review here]

The tidyverse

What is the tidyverse???


[tidyverse webpage](#)

R packages made for data science



Whole Game

Pieces of the “whole game” of data science



Whole Game

Today's plan:

- Introduce each of these principles
- Learn about `tidyverse` syntax
- Cover some examples of `tidyverse` functions that fall into each category
- Conduct our first “full analysis” by following the tidy workflow

Whole Game

Import → Tidy → Transform → Visualize → Model → Communicate

Import

Bringing our data into R

We've done this with the `read()` function family and / or the `rio` package

...Neither of those are `tidyverse` packages

But that's fine! This step is the simplest

Import: Today's dataset

Dataset on pulse rates before and after exercise

Can import from the internet

```
1 pulse_data <- rio::import("http://www.statsci.org/data/oz/ms212.txt")
```

Or download the .rds file from Canvas

```
1 pulse_data <- rio::import(here::here("labs", "Lab 5 - Intermediate R", "pulserates.rds"))
```

- .rds is a file extension for R data

Import: Today's dataset

Variable	Description
Height	Height (cm)
Weight	Weight (kg)
Age	Age (years)
Gender	Sex (1 = male, 2 = female)
Smokes	Regular smoker? (1 = yes, 2 = no)
Alcohol	Regular drinker? (1 = yes, 2 = no)
Exercise	Frequency of exercise (1 = high, 2 = moderate, 3 = low)
Ran	Whether the student ran or sat between the measurements (1 = ran, 2 = sat)
Pulse1	First pulse measurement (rate per minute)
Pulse2	Second pulse measurement (rate per minute)
Year	Year of class (93–98)

Import: Today's dataset

Let's check out our data

R Code

↺ Start Over

Run Code

```
1 # let's look at our data!
```

Whole Game

Import → **Tidy** → Transform → Visualize → Model → Communicate

Tidy

Refers to *tidy data*

“Happy families are all alike; every unhappy family is unhappy in its own way.”
— Leo Tolstoy

“Tidy datasets are all alike, but every messy dataset is messy in its own way.”
— Hadley Wickham

Tidy data follow a consistent set of organizational principles

More work up front but **much** easier to work with after

Tidy

PersonId Height Weight Year					PersonId Year Variable Value					PersonId HeightToWeight Year			
1	1	173	57	93	1	1	93	Height	173	1	1	173/57	93
2	2	179	58	93	2	1	93	Weight	57	2	2	179/58	93
3	3	167	62	93	3	2	93	Height	179	3	3	167/62	93
4	4	195	84	93	4	2	93	Weight	58	4	4	195/84	93
5	5	173	64	93	5	3	93	Height	167	5	5	173/64	93
6	6	184	74	93	6	3	93	Weight	62	6	6	184/74	93

All of these dataframes contain the same information, but one of them is much easier to work with...

Why?

Tidy: Principles

Three rules to a tidy dataset

1. Each variable is a column; each column is a variable.
2. Each observation is a row; each row is an observation.
3. Each value is a cell; each cell is a single value.

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280425583

variables

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280425583

observations

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280425583

values

Tidy: Why tidy?

1. Consistency! Any sort of consistency in your data management will be beneficial.

What's often the primary argument of functions we've used?

2. R loves vectors! Having variables in columns plays well with many functions

Tidy: Is our data tidy?

```
1 head(pulse_data, n = 10)
```

	Height	Weight	Age	Gender	Smokes	Alcohol	Exercise	Ran	Pulse1	Pulse2	Year
1	173	57	18	2	2	1	2	2	86	88	93
2	179	58	19	2	2	1	2	1	82	150	93
3	167	62	18	2	2	1	1	1	96	176	93
4	195	84	18	1	2	1	1	2	71	73	93
5	173	64	18	2	2	1	3	2	90	88	93
6	184	74	22	1	2	1	3	1	78	141	93
7	162	57	20	2	2	1	2	2	68	72	93
8	169	55	18	2	2	1	2	2	71	77	93
9	164	56	19	2	2	1	1	2	68	68	93
10	168	60	23	1	2	1	2	1	88	150	93

1. Each variable is a column; each column is a variable.
2. Each observation is a row; each row is an observation.
3. Each value is a cell; each cell is a single value.

Tidy: For later...

We will learn ways to make untidy data tidy in later classes

For now, we are good, and can proceed with our analysis!

Whole Game

Import → Tidy → Transform → Visualize → Model → Communicate

Transform

Usually, our data doesn't come to us with the variables exactly as we want them

Sometimes we need...

- New variables, based on our existing data or otherwise
- A summarized version of our dataset
- Reordered variables
- Etc., etc., etc. Any others?

Transform: **tidyverse** syntax

The **tidyverse** website mentioned that each package shares a common language for their functions

Based off of **dplyr** (one of the packages) verbs

1. The first argument is always a data frame.
2. The subsequent arguments typically describe which columns to operate on using the variable names (without quotes).
3. The output is always a new data frame.

Transform: Load the `tidyverse`

Let's get ready to actually use some `tidyverse` code!

```
1 install.packages("tidyverse")
```

```
1 library(tidyverse)
```

Note

Note the `Conflicts` section! Remember when we talked about overwriting?

Transform: More **tidyverse** syntax

One more thing...

tidyverse verbs work best in a **pipe**: `|>`

How we've been using functions:

```
output <- function(input)
```

Functions in a pipe:

```
output <- input |> function()
```

Why is this helpful?

Many functions together:

```
anotherfunction(otherfunction(function(input)))
```

Functions in a pipe:

```
input |>  
  function() |>  
  otherfunction() |>  
  another function()
```


Transform: Goals for our data

In what ways should we transform our data?

First, we need to figure out our questions!

Transform: Goals for our data

Variable	Description
Height	Height (cm)
Weight	Weight (kg)
Age	Age (years)
Gender	Sex (1 = male, 2 = female)
Smokes	Regular smoker? (1 = yes, 2 = no)
Alcohol	Regular drinker? (1 = yes, 2 = no)
Exercise	Frequency of exercise (1 = high, 2 = moderate, 3 = low)
Ran	Whether the student ran or sat between the measurements (1 = ran, 2 = sat)
Pulse1	First pulse measurement (rate per minute)
Pulse2	Second pulse measurement (rate per minute)
Year	Year of class (93–98)

Transform: Goals for our data

My questions:

1. Is the difference in pulse between the first and second measurement points different for people who ran or not?
2. Does that difference vary by whether or not someone is a regular smoker?
3. What is the relationship between first and second pulse measurements? Does this differ by smoking status?
4. Class question!! (time pending)

Transform: Gameplan

Transformation goals:

1. Make a variable with the difference between the first and second pulse measurement points
2. Code some of our numeric variables into their categories
3. Subset our data to only include the relevant columns
4. Summarize information based on our questions

Transform: Some relevant functions

<code>mutate()</code>	adds, changes, renames columns
<code>ifelse()</code>	changes a vector conditionally
<code>select()</code>	selects specified columns
<code>filter()</code>	selects specified rows
<code>summarize()</code>	summarizes data

Switching to RStudio...

Transform: Sample code

Goal 1, 2, and 3:

1. Make a variable with the difference between the first and second pulse measurement points
2. Code some of our numeric variables into their categories
3. Subset our data to only include the relevant columns

```
1 pulse_subset <- pulse_data |>
2   mutate(PulseDiff = Pulse2 - Pulse1,
3          Smokes = ifelse(Smokes == 1, "Yes",
4                          ifelse(Smokes == 2, "No", Smokes)),
5          Ran = ifelse(Ran == 1, "Ran",
6                       ifelse(Ran == 2, "Sat", Ran))) |>
7   select(Smokes, Ran, PulseDiff, Pulse1, Pulse2)
```

Transform: Sample code

Goal 4:

4. Summarize information based on our questions

```
1 pulse_subset |>  
2   summarize(.by = Ran,  
3             PulseDiffMean = mean(PulseDiff, na.rm=T))
```

Ran PulseDiffMean

```
1 Sat      -1.0000  
2 Ran      51.3913
```

```
1 pulse_subset |>  
2   summarize(.by = c(Ran, Smokes),  
3             PulseDiffMean = mean(PulseDiff, na.rm=T))
```

Ran Smokes PulseDiffMean

```
1 Sat      No      -0.8363636  
2 Ran      No      51.6511628  
3 Sat      Yes     -2.1250000  
4 Ran      Yes     47.6666667
```

Whole Game

Import → Tidy → Transform → **Visualize** → Model → Communicate

Visualize

So many numbers! Can we please visualize our data?

“The simple graph has brought more information to the data analyst’s mind than any other device.” — John Tukey


Yes! With the `ggplot2` package

A bit of a different plotting system than we’ve done before

But much more expansive and flexible (and much nicer looking, if you ask me!)

Visualize: `ggplot2`

Let's briefly explore a different question: what is the relationship between the pulse measurements at the two different time points?




Let's begin at the end:

Visualize: `ggplot2`

Let's briefly explore a different question: what is the relationship between the pulse measurements at the two different time points?

```
1 ggplot(data = pulse_subset)
```



When we just call the `ggplot()` function with our data as the first argument, we get this blank area


`ggplot()` works by adding layers and layers onto a plot object

This is not a very exciting plot, but you can think of it like an empty canvas you'll paint the remaining layers of your plot onto. — R4DS

Visualize: `ggplot2`

Let's briefly explore a different question: what is the relationship between the pulse measurements at the two different time points?

```
1 ggplot(data = pulse_subset,  
2       mapping = aes(x = Pulse1, y = Pulse2))
```



Next, we want to tell `ggplot()` what information we want in our plot and where we want it

We do this with the `aes()` (“aesthetics”) function in the `mapping =` argument

How variables are *mapped* onto specific parts of the plot

To investigate the relationship between the first and second pulse measurements, we can choose to put `Pulse1` on the x-axis and `Pulse2` on the y-axis


What do we get?

Labels, axes, and a grid

Visualize: `ggplot2`

Let's briefly explore a different question: what is the relationship between the pulse measurements at the two different time points?

```
1 ggplot(data = pulse_subset,  
2       mapping = aes(x = Pulse1, y = Pulse2)) +  
3   geom_point()
```



Still no data, though

Need to define some **geom**, a geometric display of our data

Will usually start with `geom_`

<code>geom_line()</code>	line graph
<code>geom_bar()</code>	bar graph
<code>geom_boxplot()</code>	boxplot
<code>geom_point()</code>	scatterplot

Can start addressing questions here!


Note

Take note of the missing values warning!

Visualize: `ggplot2`

Let's briefly explore a different question: what is the relationship between the pulse measurements at the two different time points?

```
1 ggplot(data = pulse_subset,  
2       mapping = aes(x = Pulse1, y = Pulse2)) +  
3   geom_point() +  
4   coord_cartesian(xlim = c(40, 180), ylim = c(40, 180))
```



These variables are conceptually similar and should probably have the same scale


Automatically generated axes are a good place to start

But good practice to consider them each time

Visualize: ggplot2

Let's briefly explore a different question: what is the relationship between the pulse measurements at the two different time points?

```
1 ggplot(data = pulse_subset,  
2       mapping = aes(x = Pulse1, y = Pulse2, color = Ran)) +  
3   geom_point()
```



Adding aesthetics and layers

Who in this group are the people who ran?


Start: do we need to modify the aesthetic or the geom?

What can we tell from this?

Visualize: ggplot2

Let's briefly explore a different question: what is the relationship between the pulse measurements at the two different time points?

```
1 ggplot(data = pulse_subset,  
2       mapping = aes(x = Pulse1, y = Pulse2, color = Ran)) +  
3   geom_point() +  
4   geom_smooth(method="lm")
```



Adding aesthetics and layers

Let's say we want to add a line to better represent the overall trend


Start: do we need to modify the aesthetic or the geom?

What can we tell from this?

Visualize: ggplot2

Let's briefly explore a different question: what is the relationship between the pulse measurements at the two different time points?

```
1 ggplot(data = pulse_subset,  
2       mapping = aes(x = Pulse1, y = Pulse2, color = Ran)) +  
3   geom_point() +  
4   geom_smooth(method="lm")
```



Adding aesthetics and layers

Let's say we want to add a line to better represent the overall trend

Start: do we need to modify the aesthetic or the geom?

What can we tell from this?


Visualize: ggplot2

Let's briefly explore a different question: what is the relationship between the pulse measurements at the two different time points?

```
1 ggplot(data = pulse_subset,  
2       mapping = aes(x = Pulse1, y = Pulse2, color = Ran)) +  
3   geom_point() +  
4   geom_smooth(method="lm") +  
5   facet_wrap(~ Smokes)
```

Let's say we want to look at this for smokers vs non-smokers

What can we tell from this?




Visualize: ggplot2

Let's briefly explore a different question: what is the relationship between the pulse measurements at the two different time points?

```
1 ggplot(data = pulse_subset,  
2       mapping = aes(x = Pulse1, y = Pulse2, color = Ran)) +  
3   geom_point() +  
4   geom_smooth(method="lm") +  
5   facet_wrap(~ Smokes) +  
6   theme_classic()
```

Let's say we want this plot to look nicer!!

Built in themes




Visualize: `ggplot2`

Let's briefly explore a different question: what is the relationship between the pulse measurements at the two different time points?

```

1 pulse_subset |>
2   ggplot(aes(x = Pulse1, y = Pulse2, color = Ran, fill = Ran)) +
3   geom_point(size = 2.4) +
4   geom_smooth(method="lm") +
5   coord_cartesian(xlim = c(40, 180), ylim = c(40, 180)) +
6   labs(x = "First Pulse", y = "Second Pulse") +
7   facet_wrap(~ Smokes, labeller = "label_both") +
8   theme_classic(base_size = 16) +
9   theme(legend.position = "bottom",
10        aspect.ratio = 1)

```



Let's say we want this plot to look nicer!!

Additional theming with `theme()`, `labs()`, and various arguments

Whole Game

Import → Tidy → Transform → Visualize → **Model** → Communicate

Model

At this point, we could take the statistical models we know and apply them to the data

We haven't learned many yet, but we have learned some!

`mean()`, `sd()` could combine to model the distribution of your data, etc.

We will do more modeling later!

Whole Game

Import → Tidy → Transform → Visualize → Model → **Communicate**

Communicate

We...have already learned ways to do this with RMarkdown!

Statistics and visualizations are great, but to tell a good story it is often useful to accompany these with text explaining the key features or importance.


And we've been practicing this all along by describing the plots!

Communicate: Pulse Measurements

```

1 pulse_subset |>
2   ggplot(aes(x = Pulse1, y = Pulse2, color = Ran, fill = Ran)) +
3   geom_point(size = 2.4) +
4   geom_smooth(method="lm") +
5   coord_cartesian(xlim = c(40, 180), ylim = c(40, 180)) +
6   labs(x = "First Pulse", y = "Second Pulse") +
7   facet_wrap(~ Smokes, labeller = "label_both") +
8   theme_classic(base_size = 16) +
9   theme(legend.position = "bottom",
10         aspect.ratio = 1)

```



In this visualization, we can see a strong positive relationship between participants' pulse measurements at the first time point and their pulse

Assignment 5