

Confidence Intervals

PSYC 2020-A01 / PSYC 6022-A01 | 2025-10-17 | Lab 9

Jessica Helmer

Outline

- Assignment 8 Review
- Confidence Intervals
- One-Sample t -test

Learning objectives:

R: CI and t -statistics in R

Assignment 7 Review

Two ways to approach two-tailed p-values

Conceptually: split your alpha onto either side of the distribution

Technically test each side (both positive and negative cutoff), but know that e.g., negative z-value isn't going to beat the positive cutoff

Compare against split alpha (2.5%, $p = .025$) on either side.

```
1 pnorm(-1.7)
```

```
[1] 0.04456546
```

Fine for just reject / retain decisions (if compare it against correct alpha)

To report the p -value specifically for the entire test, would need to multiply by 2.




```
1 pnorm(-1.7)*2
```

```
[1] 0.08913093
```

Account for probability of being more extreme in both directions

Assignment 8 Review

Remember to check the correct tail of the distribution



Confidence Intervals

Most common: 95% CI

- Interpretation: If you were to take 100 samples, 95 CIs of your 100 samples will contain the true mean

For standard normal,


CI	z cutoff	generally
99.7%	$[-3, 3]$	$[\bar{x} - 3 * SD, \bar{x} + 3 * SD]$
95%	$[-1.96, 1.96]$	$[\bar{x} - 1.96 * SD, \bar{x} + 1.96 * SD]$
68%	$[-1, 1]$	$[\bar{x} - 1 * SD, \bar{x} + 1 * SD]$

Confidence Interval and Width

Confidence Level ($1 - \alpha$) can communicate uncertainty about your results

- Designated proportion of such intervals that will include the true population value
- $\alpha = 0.01$ for 99% CI
- $\alpha = 0.05$ for 95% CI


Confidence level
proportional to
confidence interval
width



CI for Symmetric Distributions

When a distribution is symmetric, CIs for that distribution are also symmetric

- Includes normal (e.g., z-) distribution
- t distribution



Cutoff Z-Values Example

Say you want to find the cutoff z values for some confidence interval

Two examples: 95% CI and 97.3% CI (one typical, one as exercise)

We want the $\alpha = 1 - \text{confidence level}$ piece to be equal on both sides

95% CI

$\alpha = 1 - .95 = 5\%$ on both sides

So $.05/2 = .025$ on each side

Need value for .025 and $.95 + .025$
([.025, .975])

```
1 qnorm(c(.025, .975))
```

```
[1] -1.959964 1.959964
```

97.3% CI

$\alpha = 1 - .973 = 2.7\%$ on both sides

So $.027/2 = .0135$ on each side

Need value for .0135 and $.95 + .0135$
([.0135, .9865])

```
1 qnorm(c(.0135, 0.9865))
```

```
[1] -2.211518 2.211518
```

Which one has a higher confidence level? Which one has a larger width?

Confidence Interval Generally

Derived with sample mean (\bar{x}) and standard error ($\frac{\sigma}{\sqrt{n}}$)

$$CI = \bar{x} \pm z \frac{\sigma}{\sqrt{n}} \text{ or}$$

$$CI = \left[\bar{x} - z \frac{\sigma}{\sqrt{n}}, \bar{x} + z \frac{\sigma}{\sqrt{n}} \right]$$

Food for Thought

With this formula, we would only use the positive version of the z cutoff, so that the lower bound ends up lower than the mean and the higher bound ends up higher. You can also think instead as both adding the z cutoff—it just ends up becoming a minus sign because the lower bound has a negative z cutoff.

Confidence Interval Example

Let's find a 95% confidence interval of the mean for iris Petal Length

```
1 head(iris$Petal.Length)
```

```
[1] 1.4 1.4 1.3 1.5 1.4 1.7
```

```
1 x_bar <- mean(iris$Petal.Length)
```

```
2 x_sd <- sd(iris$Petal.Length)
```

```
3 n <- length(iris$Petal.Length) # although remember to be thinking about missing data
```

```
4
```

```
5 z_cutoff <- qnorm(.975)
```

```
6
```

```
7 Petal.Length.CI <- c(x_bar - z_cutoff * x_sd / sqrt(n), x_bar + z_cutoff * x_sd / sqrt(n))
```

```
8 Petal.Length.CI
```

```
[1] 3.475499 4.040501
```

If we collected samples of petal length many times, we would expect the interval [3.47, 4.04] to contain the true population mean of petal length 95% of the time.

Confidence Interval Example

Let's find a 80% confidence interval of the mean for iris Sepal Length

```
1 head(iris$Sepal.Length)
```

```
[1] 5.1 4.9 4.7 4.6 5.0 5.4
```

```
1 Sepal.Length.CI <- c(mean(iris$Petal.Length) - qnorm(.9) * sd(iris$Petal.Length) /  
2                      sqrt(sum(!is.na(iris$Sepal.Length))),  
3  
4                      mean(iris$Petal.Length) + qnorm(.9) * sd(iris$Petal.Length) /  
5                      sqrt(sum(!is.na(iris$Sepal.Length))))  
6 Sepal.Length.CI
```

```
[1] 3.573282 3.942718
```

If we collected samples of petal length many times, we would expect the interval [3.57, 3.94] to contain the true population mean of petal length 95% of the time.

Confidence Intervals and NHST

If a 95% confidence interval does not contain a value, that is mathematically equivalent to it being “significantly different” from that value.

E.g., if your null hypothesis H_0 was that the mean of petal length is no different from an expected population mean of 3.3, would you reject or retain the null hypothesis?

```
1 Petal.Length.CI
```

```
[1] 3.475499 4.040501
```

We would reject the null hypothesis because the 95% CI does not include 3.3.

**From Z to T : No longer
normal**


Z- vs. T-Distribution

T-distribution has thicker tails


As df increases, it looks more like a standard normal distribution

With $df = \infty$, exactly follows a normal distribution (so approximates with large df)


Z- vs. *T*-Distribution




Z- vs. T-Distribution




Z- vs. T-Distribution



Z- vs. T-Distribution



Z- vs. T-Distribution



***T*-Test: How many tails?**

Need to consider whether to use a “one-tail” or “two-tail” *t*-test.

One-Tail (One-Sided)

We want to test whether something is lower or higher than a value, but **not both**

Only one limit

Two-Tails (Two-Sided)


We want to test whether something is either lower or higher than a value

Two limits


Testing: How many tails?

Need to consider whether to use a “one-tail” or “two-tail” t -test.

One-Tail (One-Sided)




Two-Tails (Two-Sided)



Testing: How many tails?

Notice that two-tailed tests are harder to “beat.”



Cutoff T -Values Example

Say you want to find the cutoff t values: $t(df = 3)$ and $t(df = 37)$ for $\alpha = .05$, want both upper one-tailed and two-tailed

Which will have the larger magnitude cutoff values?

One-Tailed

All of our $\alpha = .05$ goes on the upper side

Need a cutoff for .95

```
1 qt(p = .95, df = 3, lower.tail = F)
```

```
[1] -2.353363
```

```
1 qt(p = .95, df = 37, lower.tail = F)
```

```
[1] -1.687094
```

$$t_{crit}(3) = -2.35$$

$$t_{crit}(37) = -1.69$$

Two-Tailed

$\alpha = 1 - .95 = 5\%$ on both sides

So $.05/2 = .025$ on each side

Need value for .025 and $.95 + .025$ ([.025, .975])

```
1 qt(p = c(.025, .975), df = 3)
```

```
[1] -3.182446  3.182446
```

```
1 qt(p = c(.025, .975), df = 37)
```

```
[1] -2.026192  2.026192
```

$$t_{crit}(3) = [-3.18, 3.18]$$

$$t_{crit}(37) = [-2.03, 2.06]$$

One-Sample *T*-Test Generally

Asks “is there a question between our sample and the population?”

Derived with sample mean (\bar{x}), population mean (μ), and standard error ($\frac{\bar{\sigma}}{\sqrt{n}}$)

$$t = \frac{\bar{x} - \mu}{\frac{\bar{\sigma}}{\sqrt{n}}}$$

With a *t*-test, we don't have a known population SD (σ), so we use the SD we observe in our sample $\bar{\sigma}$

Get our *t*-statistic and compare it to a critical *t* cutoff value

T-Test Example

Let's say a researcher claims the average highway miles per gallon across all cars is 30mpg. They collect a sample of 234 cars and would like you to test this. We do not know the population standard deviation.

One- or two-tailed?

```
1 head(mpg$hwy)
```

```
[1] 29 29 31 30 26 26
```

```
1 x_bar <- mean(mpg$hwy)
2 x_sd <- sd(mpg$hwy)
3 n <- length(mpg$hwy) # although remember to be thinking about missing data
4 df <- n - 1
5
6 t_cutoff <- qt(.975, df)
7
8 hwy_t_stat <- (x_bar - 30) / (x_sd / sqrt(n))
9 hwy_t_stat
```

```
[1] -16.85174
```

Our observed t -statistic exceeds our cutoff t -statistic, so we reject the null.

T-Test Function

Alternatively, we can use `t.test(x)`

- `x` = vector of numeric data
- `mu` = hypothesized population mean (default is 0)
- `alternative` = one of "two.sided", "less", "greater" (default is "two.sided")

```
t.test(mpg$hwy, mu = 30, alternative = "two.sided")
```

One Sample t-test

```
data: mpg$hwy
t = -16.852, df = 233, p-value < 2.2e-16
alternative hypothesis: true mean is not equal
to 30
95 percent confidence interval:
 22.67324 24.20710
sample estimates:
mean of x
 23.44017
```

We see a significant difference between our observed sample mean and our hypothesized population mean.

T-Test Example

Let's say a different researcher claims the average **city** miles per gallon across all cars is 30mpg. They collect a sample of 234 cars. You are confident they are wrong—you think it is certainly less than that. We do not know the population standard deviation.

One- or two-tailed?

```
1 t.test(mpg$cty, mu = 30, alternative = "less")
```

One Sample t-test

```
data: mpg$cty
t = -47.233, df = 233, p-value < 2.2e-16
alternative hypothesis: true mean is less than
30
95 percent confidence interval:
 -Inf 17.31843
sample estimates:
mean of x
```

We see a significant difference between our observed sample mean and our hypothesized population mean.

T-Test Example Output

```
1 hwy_ttest <- t.test(mpg$cty, mu = 30)
2 hwy_ttest
```

One Sample t-test

```
data: mpg$cty
t = -47.233, df = 233, p-value < 2.2e-16
alternative hypothesis: true mean is not equal
to 30
95 percent confidence interval:
 16.31083 17.40712
sample estimates:
mean of x
 16.85897
```

```
1 #?t.test
```

```
1 hwy_ttest$statistic
```

```
t
-47.23252
```

```
1 hwy_ttest$p.value
```

```
[1] 2.581229e-121
```

```
1 hwy_ttest$conf.int
```

```
[1] 16.31083 17.40712
attr(,"conf.level")
[1] 0.95
```

Can use these in RMarkdown documents with inline R chunks.

T-Tests and Confidence Intervals

The `t.test()` function includes a CI value! But...

```
t.test(mpg$cty, mu = 30) |>  
  pluck("conf.int") |>  
  c()
```

```
[1] 16.31083 17.40712
```

```
c(mean(mpg$cty) - qnorm(.975) * sd(mpg$cty) /  
  sqrt(sum(!is.na(mpg$cty)))),  
  
  mean(mpg$cty) + qnorm(.975) * sd(mpg$cty) /  
  sqrt(sum(!is.na(mpg$cty))))
```

```
[1] 16.31367 17.40428
```

They're not the same!


The built-in functions in R for confidence intervals in R often follow a *t*-distribution instead of a *z*-distribution. We want the *z*-distribution in our class, but as we saw, the differences will be very small for large samples.

T-Tests and NHST

Remember, even when the statistic is small, for two-tailed tests (because negative), we reject when we exceed the bounds of our critical value.

For one-tailed tests, it needs to exceed the bound of that tail's cutoff.

Reject or retain null?



Cohen's D

Standardized mean difference, interpretable effect size

Derived with two means (represented as μ_1 and μ_2) and standard error ($\frac{\bar{\sigma}}{\sqrt{n}}$)

$$d = \frac{\mu_1 - \mu_2}{\sigma}$$

With a t -test, we don't have a known population SD (σ), so we use the SD we observe in our sample $\bar{\sigma}$.

For a one-sample t test, μ_1 would be our sample mean \bar{x} , μ_2 would be our population mean μ , and the standard deviation σ would be our observed standard deviation $\bar{\sigma}$.

Not a statistical test—just a measure of difference.

Cohen's D Example

Let's say a researcher claims the average highway miles per gallon across all cars is 30mpg. We found earlier that the sample had a significant difference between the hypothesized and sample mean. What is the effect size of that difference?

$$d = \frac{\bar{x} - \mu}{\bar{\sigma}}$$

```
1 x_bar <- mean(mpg$hwy)
2 x_sd <- sd(mpg$hwy)
3
4 d <- (x_bar - 30) / x_sd
5 d
```

```
[1] -1.101633
```

How would be characterize this difference? What is the direction?

Cohen's D Strength Visualization

Based on mean ($\bar{x} = \text{round}(x_bar, 2) = 23.44$) and SD ($\bar{\sigma} = \text{round}(x_sd, 2) = 5.95$) of cars' highway MPG (`hwy`).

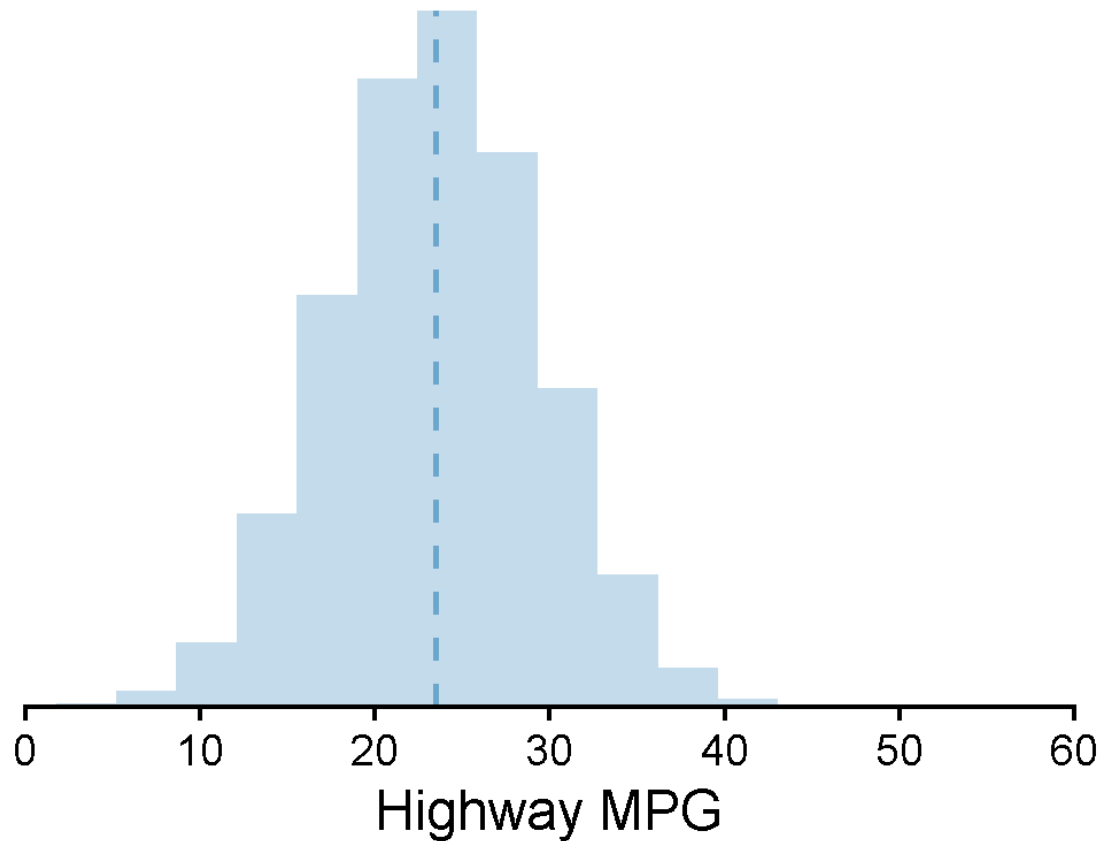
Solved backwards for (μ) based on different potential levels of d . “What would the hypothesized mean have to have been to get different levels of Cohen's d ?”

$$\bar{x} - d \times \bar{\sigma} = \mu$$

```
1 dat <- data.frame(x = rnorm(100000, x_bar, x_sd))
2
3 mu_ex1 <- x_bar - .2 * x_sd
4 mu_ex2 <- x_bar - .5 * x_sd
5 mu_ex3 <- x_bar - .8 * x_sd
```


Cohen's D Strength Visualization

“What would the hypothesized mean have to have been to get different levels of Cohen's d ?”



Cohen's D Strength Visualization

“What would the hypothesized mean have to have been to get different levels of Cohen's d ?”



Cohen's D in R

There are functions for Cohen's d in R, but they apply more to two samples of means.

- Instead of one sample and a population mean

For now, simpler to just program it ourselves $(\bar{x} - \mu) / s_x$

Two samples will come next week!

Assignment 9