

# Chi-Squared Test

PSYC 2020-A01 / PSYC 6022-A01 | 2025-11-07 | Lab 12

Jessica Helmer

# Outline

- Chi-Squared ( $\chi^2$ ) Test

Learning objectives:

R:  $\chi^2$  Tests

# Chi-Squared ( $\chi^2$ ) Test

Testing inferences about proportions

- Involve categorical variables

Comparing observed frequencies (counts) to some expected (null hypothesis) frequencies

**$\chi^2$  Test of Goodness of Fit:** expected frequencies are set by researcher (e.g., null hypothesis is equal frequencies across all groups)

**$\chi^2$  Test of Independence:** expected frequencies are those implied by independence

# Chi-Squared ( $\chi^2$ ) Test

Instead of a continuous variable, our outcome is a **count**

## Continuous

- Height
- Petal length
- Test score

## Counts

- Coin flip being heads
- Number of pets
- Goals scored

Only integers

# $\chi^2$ Test of Goodness of Fit

# $\chi^2$ Test of Goodness of Fit

Do the observed frequencies of a categorical variable differ from what expected *a priori*?

Typically, this expectation is equal occurrence across groups (but doesn't have to be).

If equal occurrence, what would be our expected frequencies?

100 coin flips

Heads	Tails
50	50

Favorite primary color of 33 students

Red	Blue	Yellow
11	11	11

Can also hypothesize proportions and convert to frequencies once you have a sample size.

# $\chi^2$ Test of Goodness of Fit: Hypotheses

$H_0$ : Observed data match the expected frequencies for the population

$H_1$ : Observed data do not match the expected frequencies for the population

- Observed frequency is significantly different than expected

$H_0: \pi_j = \pi_{j_0}$  for all categories  $j$  (i.e., difference for all categories is 0),  
where

$\pi_j$  = observed proportion

$\pi_{j_0}$  = expected proportion

$H_1: \pi_j \neq \pi_{j_0}$  for any category  $j$

# $\chi^2$ Test of Goodness of Fit Generally

$$\chi^2 = \sum_{j=1}^J \frac{(O_j - N\pi_{j_0})^2}{N\pi_{j_0}} = \sum_{j=1}^J \frac{(O_j - E_j)^2}{E_j}$$

where

$O_j$  = observed frequency

$N$  = total sample size

$E_j$  = expected frequency

$j$  = individual category out of  $J$  total categories

*The sum of the squared differences between the observed and expected frequencies, divided by the expected frequency, for each group.*



# $\chi^2$ Test of Goodness of Fit Generally

$df = J - 1$ , where


$J$  = number of categories

Use to identify critical  $\chi^2$  value from  $\chi^2$  table, R, etc.


For  $\chi^2$ , as df increases, so does the critical- $\chi^2$  value (at same  $\alpha$ )

Only ever a one-tailed test! For the upper end of the distribution.


# $\chi^2$ Distribution




# $\chi^2$ Distribution




# $\chi^2$ Distribution




# $\chi^2$ Distribution



# $\chi^2$ Distribution



# $\chi^2$ Distribution



# $\chi^2$ Test of Goodness of Fit Cutoffs

Cutoff for test of heads vs. tails

$$J = 2$$

$$\text{df} = J - 1 = 1$$

```
1 qchisq(.95, 1)
```

```
[1] 3.841459
```

```
1 qchisq(.05, 1, lower.tail = F)
```

```
[1] 3.841459
```

Cutoff for test of favorite color in the rainbow

$$J = 7 \text{ (ROYGBIV)}$$

$$\text{df} = J - 1 = 6$$

```
1 qchisq(.95, 6)
```

```
[1] 12.59159
```

```
1 qchisq(.05, 6, lower.tail = F)
```

```
[1] 12.59159
```

# $\chi^2$ Test of Goodness of Fit Example

You are an education psychologist interested in college students' choice of majors. You've grouped majors into five categories: STEM, Social Sciences, Humanities, Arts, and Business. You'd like to test whether there are differences in the population of students choosing those majors. You sample 100 students and ask what major they chose.

First, what are your expected frequencies for each category?

	<b>STEM</b>	<b>Social Sciences</b>	<b>Humanities</b>	<b>Arts</b>	<b>Business</b>
Expected					
Observed					

# $\chi^2$ Test of Goodness of Fit Example

You are an education psychologist interested in college students' choice of majors. You've grouped majors into five categories: STEM, Social Sciences, Humanities, Arts, and Business. You'd like to test whether there are differences in the population of students choosing those majors. You sample 100 students and ask what major they chose.

First, what are your expected frequencies for each category?

	<b>STEM</b>	<b>Social Sciences</b>	<b>Humanities</b>	<b>Arts</b>	<b>Business</b>
Expected	20	20	20	20	20
Observed					

# $\chi^2$ Test of Goodness of Fit Example

You are an education psychologist interested in college students' choice of majors. You've grouped majors into five categories: STEM, Social Sciences, Humanities, Arts, and Business. You'd like to test whether there are differences in the population of students choosing those majors. You sample 100 students and ask what major they chose.

Next, what is your  $\chi^2$  cutoff value?

	<b>STEM</b>	<b>Social Sciences</b>	<b>Humanities</b>	<b>Arts</b>	<b>Business</b>
Expected	20	20	20	20	20
Observed					

# $\chi^2$ Test of Goodness of Fit Example

## Cutoff in R

```
1 df <- 5 - 1  
2 chisq_crit <- qchisq(.95, df)  
3 chisq_crit
```

```
[1] 9.487729
```

# $\chi^2$ Test of Goodness of Fit Example

You are an education psychologist interested in college students' choice of majors. You've grouped majors into five categories: STEM, Social Sciences, Humanities, Arts, and Business. You'd like to test whether there are differences in the population of students choosing those majors. You sample 100 students and ask what major they chose.

Here are the observed values:

	STEM	Social Sciences	Humanities	Arts	Business
Expected	20	20	20	20	20
Observed	65	10	5	5	15

# $\chi^2$ Test of Goodness of Fit Example

$$\chi^2 = \sum_{j=0}^J \frac{(O_j - N\pi_j)^2}{N\pi_j} = \sum_{j=0}^J \frac{(O_j - E_j)^2}{E_j}$$

```
1 obs_chisq <- (65 - 20)^2 / 20 +
2   (10 - 20)^2 / 20 +
3   (5 - 20)^2 / 20 +
4   (5 - 20)^2 / 20 +
5   (15 - 20)^2 / 20
6 obs_chisq
```

```
[1] 130
```

```
1 obs_chisq > chisq_crit
```

```
[1] TRUE
```

We can reject the null in favor of the alternative that the number of students is not the same across majors.

# $\chi^2$ Test of Goodness of Fit R Function

```
majordat <- data.frame(major = c("STEM", "Social.Sciences", "Humanities", "Arts", "Business"),  
observed = c(65, 10, 5, 5, 15))  
  
chisq <- majordat |>  
  select(observed) |>  
  chisq.test()  
  
chisq
```

Chi-squared test for given probabilities

```
data: select(majordat, observed)  
X-squared = 130, df = 4, p-value < 2.2e-16
```

1 chisq\$expected

```
[1] 20 20 20 20 20
```

1 chisq\$statistic

X-squared  
130

1 chisq\$p.value

```
[1] 3.89406e-27
```

# $\chi^2$ Test of Goodness of Fit Example

Let's restart and say that you were a educational psychologist interested in how students in STEM schools choose their major. In this situation, you had expected more students to select a STEM major than other majors.

These are your new expected frequencies (60% STEM, evenly divided across the rest):

	<b>STEM</b>	<b>Social Sciences</b>	<b>Humanities</b>	<b>Arts</b>	<b>Business</b>
Expected	60	10	10	10	10
Observed	65	10	5	5	15

# $\chi^2$ Test of Goodness of Fit R Function

```
chisq.test(x = c(65, 10, 5, 5, 15),  
           p = c(60, 10, 10, 10, 10),  
           rescale.p = T)
```

Chi-squared test for given probabilities

```
data: c(65, 10, 5, 5, 15)  
X-squared = 7.9167, df = 4, p-value = 0.09468
```

```
chisq.test(x = c(65, 10, 5, 5, 15),  
           p = c(.60, .10, .10, .10, .10))
```

Chi-squared test for given probabilities

```
data: c(65, 10, 5, 5, 15)  
X-squared = 7.9167, df = 4, p-value = 0.09468
```

- **x** = data of frequencies
- **p** = expected probabilities
- **rescale.p** = [T/F], will rescale **p** to probabilities if you prefer to input frequencies

# $\chi^2$ Test of Independence

# $\chi^2$ Test of Independence

Do the observed frequencies of a categorical variable differ from what expected via properties of independence?

- Now considering more than one type of category
- E.g., Number of pets by cat vs. dog people and undergrad vs. post-graduate

	Cat Person	Dog Person
Undergrad		
Post-Graduate		

Expected frequencies calculated from the marginal totals of the contingency table

- Expected values determined from data

# $\chi^2$ Test of Independence: Hypotheses

$H_0$ : Categorical Variable 1 and Categorical Variable 2 are independent (i.e., have no association)

$H_1$ : Categorical Variable 1 and Categorical Variable 2 are not independent (i.e., have some association)

- Do expected frequencies match those expected via independence?

Mathematically the same as before:

$H_0: \pi_j = \pi_{j_0}$  for all categories  $j$ ; difference for all categories is 0, where  
 $\pi_j$  = observed proportion

$\pi_{j_0}$  = expected proportion

$H_1: \pi_j \neq \pi_{j_0}$  for any category  $j$

Just changing the expected frequencies



# $\chi^2$ Test of Independence Generally

$$\chi^2 = \sum_{j_1=1}^{J_1} \sum_{j_2=1}^{J_2} \frac{(O_j - N\pi_{j_0})^2}{N\pi_{j_0}} = \sum_{j_1=1}^{J_1} \sum_{j_2=1}^{J_2} \frac{(O_j - E_j)^2}{E_j}$$

Basically identical formula to calculate the  $\chi^2$ .

The difference is that our expected frequencies  $E_j$  are derived from the data. Just like before, we calculate the difference from each cell, but now we have two categories that make up the cell.

*The sum of the squared differences between the observed and expected frequencies, divided by the expected frequency, for each cell.*

# $\chi^2$ Test of Independence Cutoffs

What is the critical  $\chi^2$  value in a test of independence for a 3x2?

$$df = (J_1 - 1) \times (J_2 - 1) \text{ or}$$

$$df = (n_{rows} - 1) \times (n_{cols} - 1)$$

```
1 df <- (3 - 1) * (2 - 1)
2 df
```

```
[1] 2
```

```
1 qchisq(.95, df)
```

```
[1] 5.991465
```

# $\chi^2$ Test of Independence Example

You are another educational psychologist, and you're interested in examining your department's student body makeup. You wonder whether there is an association between students' level (graduate vs. undergraduate) and whether they are an in-state student or not. You sample 50 students from your department, and these are the frequencies you observe.

	In-State	Out-of-State
Undergrad	25	10
Graduate	5	10

First, what are your expected frequencies for each category?

# $\chi^2$ Test of Independence Example: Calculating Expected Frequencies

First, get the marginal totals for each column

	In-State	Out-of-State	<i>Total</i>
<b>Undergrad</b>	25	10	
<b>Graduate</b>	5	10	
<b><i>Total</i></b>			

# $\chi^2$ Test of Independence Example: Calculating Expected Frequencies

First, get the marginal totals for each column

	In-State	Out-of-State	Total
Undergrad	25	10	35
Graduate	5	10	15
Total	30	20	50

# $\chi^2$ Test of Independence Example: Calculating Expected Frequencies

	In-State	Out-of-State	Total
Undergrad	25	10	35
Graduate	5	10	15
Total	30	20	50

	In-State	Out-of-State	Total
Undergrad	21	14	35
Graduate	9	6	15
Total	30	20	50

Expected cell frequency =  $E_j = \frac{\text{column total} \times \text{row total}}{N}$

- Undergrad & In-State =  $\frac{35 \times 30}{50} = 21$
- Undergrad & Out-of-State =  $\frac{35 \times 20}{50} = 14$
- Graduate & In-State =  $\frac{15 \times 30}{50} = 9$
- Graduate & Out-of-State =  $\frac{15 \times 20}{50} = 6$



# $\chi^2$ Test of Goodness of Fit Example

$$\chi^2 = \sum_{j=0}^J \frac{(O_j - N\pi_j)^2}{N\pi_j} = \sum_{j=0}^J \frac{(O_j - E_j)^2}{E_j}$$

```
data.frame(type = c("Undergrad", "Graduate", "Total"),
           `In-State` = c(25, 5, 30),
           `Out-of-State` = c(10, 10, 20),
           Total = c(35, 15, 50))
```

	type	In.State	Out.of.State	Total
1	Undergrad	25	10	35
2	Graduate	5	10	15
3	Total	30	20	50

```
chisq_crit <- qchisq(.95, df = (2 - 1) * (2 - 1))
chisq_crit
```

```
[1] 3.841459
```

```
obs_chisq <- (25 - 21)^2 / 21 +
(10 - 14)^2 / 14 +
(5 - 9)^2 / 9 +
(10 - 6)^2 / 6
obs_chisq
```

```
[1] 6.349206
```

```
obs_chisq > chisq_crit
```

```
[1] TRUE
```

# $\chi^2$ Test of Independence R Function

```
studdat <- data.frame(type = c("Undergrad", "Graduate"),  
                      `In-State` = c(25, 5),  
                      `Out-of-State` = c(10, 10))  
  
studdat |>  
  select(-type) |>  
  chisq.test(correct = F)
```

- `correct` = [T/F], applies a “continuity correction” to the  $\chi^2$  value. Default is T. Change to F!

Pearson's Chi-squared test

```
data: select(studdat, -type)  
X-squared = 6.3492, df = 1, p-value = 0.01174
```

We can reject the null in favor of the alternative that students' level (graduate vs. undergraduate) and whether they are an in-state student or not are not independent.

# Count Data

# Working With Count Data

```
1 dat |> head(12)
```

	student	major	type	
1	1	STEM	Non-STEM	School
2	2	Humanities	STEM	School
3	3	Humanities	STEM	School
4	4	Social Sciences	Non-STEM	School
5	5	Humanities	STEM	School
6	6	Arts	Non-STEM	School
7	7	Social Sciences	STEM	School
8	8	Social Sciences	STEM	School
9	9	STEM	STEM	School
10	10	STEM	STEM	School
11	11	Business	STEM	School
12	12	Social Sciences	STEM	School

# Working With Count Data: `table()`

Some ways to generate frequencies from categorical variables

```
1 table(dat$major)
```

Arts	Business	Humanities	Social Sciences	STEM
23	13	22	19	23

```
1 table(dat$major) |>  
2 data.frame()
```

	Var1	Freq
1	Arts	23
2	Business	13
3	Humanities	22
4	Social Sciences	19
5	STEM	23

# Working With Count Data: `table()`

```
1 table(dat$major, dat$type)
```

	Non-STEM	School	STEM	School
Arts		13		10
Business		4		9
Humanities		10		12
Social Sciences		10		9
STEM		16		7

```
1 table(dat$major, dat$type) |>  
2 data.frame()
```

	Var1	Var2	Freq	
1	Arts	Non-STEM	School	13
2	Business	Non-STEM	School	4
3	Humanities	Non-STEM	School	10
4	Social Sciences	Non-STEM	School	10
5	STEM	Non-STEM	School	16
6	Arts	STEM	School	10
7	Business	STEM	School	9
8	Humanities	STEM	School	12
9	Social Sciences	STEM	School	9
10	STEM	STEM	School	7

# Working With Count Data: `summarize()`

Some ways to generate frequencies from categorical variables

```
1 dat |>  
2   summarize(.by = c(major, type),  
3             n = n())
```

	major	type	n		
1	STEM	Non-STEM	School	16	
2	Humanities	STEM	School	12	
3	Social Sciences	Non-STEM	School	10	
4		Arts	Non-STEM	School	13
5	Social Sciences	STEM	School	9	
6		STEM	STEM	School	7
7	Business	STEM	School	9	
8	Humanities	Non-STEM	School	10	
9		Arts	STEM	School	10
10	Business	Non-STEM	School	4	

```
1 dat |>  
2   summarize(.by = c(major, type),  
3             n = n()) |>  
4   arrange(major, type)
```

	major	type	n		
1	Arts	Non-STEM	School	13	
2	Arts	STEM	School	10	
3	Business	Non-STEM	School	4	
4	Business	STEM	School	9	
5	Humanities	Non-STEM	School	10	
6	Humanities	STEM	School	12	
7		STEM	Non-STEM	School	16
8		STEM	STEM	School	7
9	Social Sciences	Non-STEM	School	10	
10	Social Sciences	STEM	School	9	

# Working With Count Data: count()

count(x, grouping\_var1, grouping\_var2, etc.)

- **x** = dataframe
- **grouping\_vars** = variable names to group by

Basically shorthand for:

```
1 df |>  
2   summarise(.by = c(a, b),  
3             n = n())
```

Also seems to **arrange** by default

# Working With Count Data: count()

count(x, grouping\_var1, grouping\_var2, etc.)

- `x` = dataframe
- `grouping_vars` = variable names to group by

```
1 dat |>  
2 summarize(.by = c(major, type),  
3             n = n())
```

	major	type	n
1	STEM	Non-STEM	School 16
2	Humanities	STEM	School 12
3	Social Sciences	Non-STEM	School 10
4		Arts	Non-STEM School 13
5	Social Sciences	STEM	School 9
6		STEM	STEM School 7
7	Business	STEM	School 9
8	Humanities	Non-STEM	School 10
9		Arts	STEM School 10
10	Business	Non-STEM	School 4

```
1 dat |>  
2 count(major, type)
```

	major	type	n
1		Arts	Non-STEM School 13
2		Arts	STEM School 10
3		Business	Non-STEM School 4
4		Business	STEM School 9
5		Humanities	Non-STEM School 10
6		Humanities	STEM School 12
7		STEM	Non-STEM School 16
8		STEM	STEM School 7
9	Social Sciences	Non-STEM	School 10
10	Social Sciences	STEM	School 9

# Working With Count Data: `summarize()`

If going to put into the `chisq.test()` function, want to pivot and get rid of additional columns.

Many ways to do this, but one example:

```
chisqdat <- dat |>  
  count(major, type) |>  
  pivot_wider(names_from = type,  
              values_from = n) |>  
  select(-major)  
chisqdat
```

```
# A tibble: 5 × 2  
  `Non-STEM School` `STEM School`  
    <int>        <int>  
1        13          10  
2         4           9  
3        10          12  
4        16           7  
5        10           9
```

```
chisq.test(chisqdat, correct = F)
```


Pearson's Chi-squared test

```
data: chisqdat  
X-squared = 5.7312, df = 4, p-value = 0.2201
```

# Working With Count Data: `geom_bar()`

Plot


Code



# Working With Count Data: `geom_bar()`

Plot

Code



# Assignment 12